

SRAV S

AI/ML Engineer

Los Angeles, CA | +1 408-618-3557 | ssrav641@gmail.com | [LinkedIn](#)

SUMMARY

AI/ML Engineer with 5+ years of experience building and deploying production AI systems and AI applications across traditional machine learning, Generative AI, LLMs, Agentic AI, and retrieval. Specializes in Agentic AI, designing multi-agent architectures, agent orchestration pipelines, RAG and context-grounding systems, and API-driven AI services using Python, LangChain, LangGraph, and vector search to support complex planning, analysis, and decision-support use cases. Strengthens AI reliability through LLM and agent evaluation, human-in-the-loop validation, adversarial and regression testing, hallucination detection, and continuous monitoring. Earlier experience delivering end-to-end predictive ML for transportation pricing, from feature engineering and model development through API integration and production monitoring. Partners closely with engineering, data, and business teams to turn complex requirements into scalable, maintainable AI solutions with measurable operational impact.

PROFESSIONAL EXPERIENCE

Scale AI

Jan 2025 - Present

AI/ML Engineer

Los Angeles, CA

- Designed and built Agentic AI and multi-agent systems in Python using LLMs for complex planning, analysis, and decision-support use cases, coordinating 5 specialized agents across 12 production workflows to run parallel analysis, critique proposed plans, and consolidate findings into a single recommendation.
- Engineered agent orchestration pipelines with LangChain and LangGraph to route tasks across agents, manage state, and coordinate 15+ integrated tools and APIs through function calling, merging agent outputs into structured, actionable decision-support results.
- Integrated LLM applications with external APIs, data services, and simulation workflows through tool-calling interfaces, enabling agents to retrieve live context and support scenario-based analysis.
- Implemented Retrieval-Augmented Generation (RAG) and context-grounding pipelines over 75K+ indexed documents using vector and semantic search, constraining LLM outputs to trusted sources and reducing hallucination rate by 12%.
- Developed production-ready AI services in Python exposed through REST APIs, embedding generative AI capabilities into modular application architectures and existing enterprise workflows.
- Designed LLM, agent, and RAG evaluation workflows including LLM-as-a-judge scoring spanning 40+ evaluation scenarios to measure response quality, reliability, consistency, and grounding, improving response-quality scores by 10% and catching regressions before they reached production.
- Built human-in-the-loop validation processes that routed AI outputs to domain experts and fed structured feedback into iterative agent improvement, cutting expert review time per output by 25%.
- Built a suite of 150+ adversarial and regression test cases to surface hallucinations, inconsistent responses, and edge-case failures before release, turning each confirmed failure into a permanent regression case so the same defect could not ship twice.
- Containerized AI services with Docker and applied LLMOps practices across CI/CD, automated testing, logging, monitoring, and AI observability to run scalable generative AI workloads reliably in production.
- Partnered with AI/ML engineers, software engineers, and domain stakeholders to translate complex requirements into scalable AI solutions, aligning agent workflows with operational and business objectives.

Infosys

Aug 2019 - Jan 2023

Machine Learning Engineer

India

- Built machine learning solutions for transportation and logistics spot-market pricing, replacing manual quoting with data-driven freight price recommendations for the pricing business.
- Analyzed historical shipment and pricing data across routes and shipment types to identify the cost drivers with the strongest influence on transportation pricing.
- Prepared and transformed structured logistics data for modeling by resolving missing values, outliers, and data inconsistencies, improving the quality and consistency of training data.
- Engineered features from shipment distance, origin-destination pairs, load characteristics, historical pricing, fuel prices, and seasonal factors, improving model predictive performance.
- Built, compared, and validated regression-based machine learning models to predict freight prices, selecting the approach that delivered the most consistent predictions on held-out historical data.
- Tuned model hyperparameters and refined input features to reduce pricing prediction error by 15% on held-out validation data.
- Developed reusable data-processing and model-training workflows and integrated predictions into downstream business systems through APIs, enabling automated pricing recommendations for new shipments.
- Monitored model outputs and data quality in production to detect unusual predictions, shifting data patterns, and performance degradation.

- Partnered with data engineers, software developers, and business stakeholders to translate pricing requirements into production ML solutions, improving pricing accuracy, increasing bid-win rates by 7%, and delivering \$300K+ in annual savings for the logistics business.

PROJECT

Multi-Agent Enterprise Decision Support System

- Designed multi-agent workflows using Python, LangChain, and LangGraph for planning, analysis, validation, and decision-support tasks with integrated tools and APIs.
- Implemented LLM evaluation, adversarial testing, regression testing, and human feedback loops to improve agent reliability and detect hallucinations.

TECHNICAL SKILLS

Programming: Python, SQL

AI/ML: Machine Learning, Deep Learning, Neural Networks, Predictive Modeling, Regression, NLP, Feature Engineering, Model Evaluation, Model Optimization

Generative AI & LLMs: Generative AI, LLMs, Transformer Models, Prompt Engineering, Context Engineering, Embeddings

Agentic AI & LLM Engineering: AI Agents, Agentic AI, Multi-Agent Systems, Agent Orchestration, LangChain, LangGraph, Tool Calling, Function Calling, MCP

RAG & Retrieval: RAG, Context Grounding, Vector Databases, Vector Search, Semantic Search, Information Retrieval, Reranking

AI Evaluation & Reliability: LLM Evaluation, Agent Evaluation, RAG Evaluation, LLM-as-a-Judge, HITL, Adversarial Testing, Regression Testing, Hallucination Detection

MLOps & AI Infrastructure: MLOps, LLMOps, ML Pipelines, Model Deployment, Model Serving, Model Monitoring, Experiment Tracking, Model Versioning, Continuous Evaluation, AI Observability

Cloud & Deployment: AWS, Azure, GCP, Docker, Kubernetes, CI/CD, Distributed Systems, Scalable AI Infrastructure

Data Engineering: Data Pipelines, ETL/ELT, Data Processing, Data Transformation, Data Quality, Data Validation, Batch Processing, Real-Time Processing

ML/AI Frameworks: PyTorch, TensorFlow, scikit-learn, Hugging Face, MLflow, PySpark, Apache Spark

GPU Computing: CUDA, GPU Computing

EDUCATION

Master of Science, Computer and Information Science
Concordia University Wisconsin

Mar 2023 – Dec 2024